

CS 60002: Distributed Systems

T10: Distributed Mutual Exclusion

Department of Computer Science
and Engineering



INDIAN INSTITUTE OF TECHNOLOGY
KHARAGPUR



Sandip Chakraborty
sandipc@cse.iitkgp.ac.in

Mutual Exclusion (Mutex)

- A program object to prevent simultaneous access to the shared resources, used in concurrent programming with a **critical section** (a part of code where processes access shared resources)
 - Well understood in case of shared memory systems
 - **Reader-Writer:** Two processes can read simultaneously, but read-write or write-write cannot happen simultaneously

Mutual Exclusion (Mutex)

- A program object to prevent simultaneous access to the shared resources, used in concurrent programming with a **critical section** (a part of code where processes access shared resources)
 - Well understood in case of shared memory systems
 - **Reader-Writer**: Two processes can read simultaneously, but read-write or write-write cannot happen simultaneously
- **Requirements:**
 - At most one process should execute the critical section (**Safety**)
 - If more than one processes request to execute the critical section, at least one of them enters (**liveness**)
 - A process executes the critical section for a finite time (**no starvation**)
 - The requests are granted in some order (**fairness**)

Distributed Mutual Exclusion

- **Permission-based Algorithms**

- A process takes permission from all other processes (or a subset of that) before entering the critical section
- **Permission from all:** Costly, good for small systems
- **Permission from subset:** scalable, widely used; however, the main challenge is how to choose the subset for each process?

Distributed Mutual Exclusion

- **Permission-based Algorithms**

- A process takes permission from all other processes (or a subset of that) before entering the critical section
- **Permission from all:** Costly, good for small systems
- **Permission from subset:** scalable, widely used; however, the main challenge is how to choose the subset for each process?

- **Token-based Algorithms**

- Single token in the system
- A process enters the critical section if it has a token
- There are different algorithms based on how the token is circulated across different requesting processes

Adaptive vs Non-adaptive Algorithms

- Less requests for entering the critical section → Low load
- Performance of a distributed mutual exclusion algorithm under low load should be better than that under high load
- An algorithm is called adaptive if the performance is dependent on the load
- But, how do we measure the performance?

Metrics for Performance Measurements

- **Message complexity** -- Number of messages per critical section entry
- **Synchronization delay** – Time required to enter the critical section after a process leaves the critical section
- **Response time** – The time interval a request waits for its critical section execution to be over after its request messages have been sent out
- **Throughput** – The rate at which critical sections are executed, computed as $1/(\text{synchronization delay} + \text{critical section execution time})$

Permission-based Algorithms

Lamport's Algorithm

- Permission from all, contention-based algorithm
 - Processes contend with each other to enter the critical section
- Uses logical clock to order the messages
 - Indeed, the algorithm was given as an example of the use of Logical clocks
- System model: Asynchronous, Completely-connected topology, Reliable and FIFO channel

Lamport's Algorithm

- Every process maintains a queue of pending requests for entering critical section in that order
 - The requests in the queue is ordered based on the logical clock timestamp
 - Total order is enforced by including process IDs in the timestamp

Lamport's Algorithm

- Every process maintains a queue of pending requests for entering critical section in that order
 - The requests in the queue is ordered based on the logical clock timestamp
 - Total order is enforced by including process IDs in the timestamp
- **Process i requests to enter the critical section**
 - Send timestamped REQUEST (ts_i, i) to all other processes
 - Put (ts_i, i) in its own queue

Lamport's Algorithm

- Every process maintains a queue of pending requests for entering critical section in that order
 - The requests in the queue is ordered based on the logical clock timestamp
 - Total order is enforced by including process IDs in the timestamp
- **Process i requests to enter the critical section**
 - Send timestamped REQUEST (ts_i, i) to ALL other processes
 - Put (ts_i, i) in its own queue
- **On receiving a request (ts_i, i)**
 - Send timestamped REPLY to the requesting node i , requests are granted in the order of timestamp
 - Put the REQUEST (ts_i, i) in the queue

Lamport's Algorithm

- **Process i enters the critical section if**
 - (ts_i, i) is at the top of its own queue
 - Process i has received a message (REQUEST or REPLY) with a timestamp larger than ts_i from ALL other processes


Lamport's Algorithm

- **Process i enters the critical section if**
 - (ts_i, i) is at the top of its own queue
 - Process i has received a message (REQUEST or REPLY) with a timestamp larger than ts_i from ALL other processes
- **Process i releases the critical section**
 - i removes the request from its own queue and sends a timestamped RELEASE message to ALL other processes


Lamport's Algorithm

- **Process i enters the critical section if**
 - (ts_i, i) is at the top of its own queue
 - Process i has received a message (REQUEST or REPLY) with a timestamp larger than ts_i from ALL other processes
- **Process i releases the critical section**
 - i removes the request from its own queue and sends a timestamped RELEASE message to ALL other nodes
- **On receiving a RELEASE message from i**
 - i 's request is removed from the local request queue


Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration




Lamport's Algorithm -- Illustration



Lamport's Algorithm -- Illustration



Lamport's Algorithm -- Illustration



Some Observations

- Timestamp of a REPLY message is greater than the timestamp of the corresponding REQUEST message
- The REPLY message from process i to process j ensures that j knows of all requests of i prior to sending the REPLY
 - This ensures that j knows any possible requests from i with a lower timestamp
 - Requires FIFO channel
- Requests are granted in the order of increasing timestamp

Performance Metrics

- $3(n-1)$ messages are required per critical section invocation
 - $(n-1)$ REQUEST messages
 - $(n-1)$ REPLY messages
 - $(n-1)$ RELEASE messages
- Synchronization delay equals to maximum message transmission time

Ricart-Agrawala Algorithm

- Improvement over Lamport's Algorithm
 - Developed by Glenn Ricart and Ashok Agrawala
 - Ricart, Glenn; Agrawala, Ashok K. (1981). "An optimal algorithm for mutual exclusion in computer networks". Communications of the ACM. 24 (1): 9–17

Ricart-Agrawala Algorithm

- Improvement over Lamport's Algorithm
 - Developed by Glenn Ricart and Ashok Agrawala
 - Ricart, Glenn; Agrawala, Ashok K. (1981). "An optimal algorithm for mutual exclusion in computer networks". Communications of the ACM. 24 (1): 9–17
- **Core Idea:**
 - Process i has a request timestamp of ts_i and process j has a request timestamp of ts_j
 - $ts_j < ts_i \rightarrow i$ cannot enter the critical section before j
 - Process j needs not send a REPLY immediately to process i in this case
 - Process j sends the REPLY after it completes its critical section
 - No separate RELEASE message is needed

Ricart-Agrawala Algorithm

- **To Request the Critical Section:**

- Send timestamped REQUEST message (ts_i, i)

- **On receiving Request (ts_i, i) at j :**

- Send REPLY to i if
 - j is neither requesting nor executing critical section or
 - if j is requesting and i 's request timestamp is smaller than j 's request timestamp.
- Otherwise, defer the request.

Ricart-Agrawala Algorithm

- **To enter critical section:**
 - i enters critical section on receiving REPLY from all nodes
- **To release critical section:**
 - Send REPLY to all deferred requests

Analysis of the Algorithm

- Requests are granted in order of increasing timestamps
- Does not require FIFO, but requires knowledge of all other processes in the network
- $2(n-1)$ messages per critical section invocation
- Synchronization delay equals to the maximum message transmission time
- Failure of a process may cause starvation – needs separate fault detection algorithm

Maekawa's Algorithm

- Permission is obtained from a subset of other processes
 - The subset is called the **Request Set** (or **Quorum Set**)
 - Separate request set R_i for each process i

Maekawa's Algorithm

- Permission is obtained from a subset of other processes
 - The subset is called the **Request Set** (or **Quorum Set**)
 - Separate request set R_i for each process i
- **Requirements for the Quorum Set:**
 - For all i, j : $R_i \cap R_j \neq \emptyset$
 - For all i : $i \in R_i$
 - For all i : $|R_i| = K$ for some K
 - Any process i is contained in exactly D quorum sets, for some D
- $K = D$
- $K \geq \sqrt{N-1}$

Maekawa's Algorithm – A Simple Version

- **To request critical section:**
 - i sends REQUEST message to all processes in its quorum set R_i
- **On receiving a REQUEST message**
 - Send a REPLY message if no REPLY message has been sent since the last RELEASE message is received.
 - Update status to indicate that a REPLY has been sent.
 - Otherwise, queue up the REQUEST

Maekawa's Algorithm – A Simple Version

- **To enter critical section:**

- i enters critical section after receiving REPLY from all processes in R_i

- **To release critical section:**

- Send RELEASE message to all processes in R_i
- On receiving a RELEASE message, send REPLY to next process in the queue and delete the process from the queue.
- If queue is empty, update status to indicate no REPLY message has been sent since last RELEASE is received.

Maekawa's Algorithm – Analysis


- The REPLY message works like a vote
 - To enter the critical section, every process needs to obtain the vote from ALL the processes in its quorum set
- **Safety:** Any two quorum set has at least one process in common
 - The common processes ensure the total ordering of requesting processes
- **Message Complexity:** $3 \times \sqrt{N}$
- **Synchronization delay** = 2 message transmission delays
- A major issue with the algorithm: **deadlock is possible**
 - $i \in R_j, j \in R_i$, i waits for the reply from j, j waits for the reply from i
 - Try to solve this problem with additional messages and logical clock timestamp (indeed, Maekawa's algorithm already does that)
- How do we build the quorum set?

Properties of a Quorum

- The set of Quorums is called a **Coterie**. The following properties must hold for quorums in a coterie.
- **Intersection:** For every quorum $Q1$ and $Q2$ in C , $Q1 \cap Q2 \neq \emptyset$.
 - For example, $\{1, 2, 3\}$, $\{3, 4, 5\}$ and $\{7, 8, 9\}$ cannot be quorums in a coterie
- **Minimality:** There should not be any two quorums $Q1$ and $Q2$ in C , such that $Q1 \subseteq Q2$.
 - For example, $\{1, 2, 3\}$ and $\{1, 2\}$ cannot be quorums of a coterie
 - **Ensure efficiency rather than correctness**


Properties of a Quorum

- Let x is process in quorum Q . If x wants to invoke mutual exclusion, the **Intersection property** ensures that, there is at least one process in the quorum Q that is common to the quorum of every other processes
 - These common processes send permission to only one processes at a time, thus ensuring mutual exclusion




Properties of a Quorum

- Let x is process in quorum Q . If x wants to invoke mutual exclusion, the **Intersection property** ensures that, there is at least one process in the quorum Q that is common to the quorum of every other processes
 - These common processes send permission to only one processes at a time, thus ensuring mutual exclusion



Properties of a Quorum

- Let x is process in quorum Q . If x wants to invoke mutual exclusion, the **Intersection property** ensures that, there is at least one process in the quorum Q that is common to the quorum of every other processes
 - These common processes send permission to only one processes at a time, thus ensuring mutual exclusion




Maekawa's Algorithm – Building the Quorum Set

- One possible solution is to use a 2-D grid
 - Need $2 \cdot \sqrt{N}$ votes and $6 \cdot \sqrt{N}$ messages

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

Agarwal-El Abbadi Algorithm

- Uses '**tree-structured**' quorums – All the processes are logically organized in a complete binary tree
- Each quorum represents a path from the root to the leaf




- Eight quorums

- {1, 2, 4, 8}
- {1, 2, 4, 9}
- {1, 2, 5, 10}
- {1, 2, 5, 11}
- {1, 3, 6, 12}
- {1, 3, 6, 13}
- {1, 3, 7, 14}
- {1, 3, 7, 15}

Agarwal-El Abbadi Algorithm

- Number of processes in a quorum: $\log n$
- Handle process failures till $\log n$
 - If any process fails, the algorithm substitutes for that process two possible paths starting from the process's two children and ending in leaf nodes.
 - Cannot form quorums when failure is more than $\log n$



- Eight quorums when **process 3 fails**
 - {1, 2, 4, 8}
 - {1, 2, 4, 9}
 - {1, 2, 5, 10}
 - {1, 2, 5, 11}
 - {1, 6, 12, 7, 14}
 - {1, 6, 13, 7, 15}
 - {1, 6, 13, 7, 14}
 - {1, 6, 12, 7, 15}

Token-based Algorithms

Broad Idea

- A single token circulates in the system; a process enters the critical section when it has the token
- Obvious mutual exclusion – no more than one process can possess the token at any point in time
- Algorithms differ in how to find and get the token
 - Token circulates, processes use it when it passes through them
 - Token stays at the process that used it last, other processes request for it when needed – need to differentiate between old and current requests

Suzuki Kasami Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it
- **Issues:**
 - Current vs Outdates requests
 - Determining processes with pending requests
 - Deciding which process to give the token to

Suzuki Kasami Algorithm

- The token (called the PRIVILEGE message in the algorithm):
 - Queue (FIFO) Q of requesting processes
 - $LN[1...n]$: Sequence number of requests that j executed most recently
- The request message:
 - $REQUEST(i, k)$: Request message from process i for its k^{th} critical section execution
- Other data structures:
 - $RN_i[1...n]$ for each node i, where $RN_i[j]$ is the largest sequence number received so far by i in a REQUEST message from j

Suzuki Kasami Algorithm

- To request critical section:
 - If i does not have the token, increment $RN_i[i]$ and send $REQUEST(i, RN_i[i])$ to all processes
 - If i has the token already, enter critical section if the token is idle (no pending requests in token Q), else follow rule to release critical section
- On receiving $REQUEST(i, sn)$ at j :
 - set $RN_j[i] = \max(RN_j[i], sn)$
 - If j has the token and the token is idle, send it to i if $RN_j[i] = LN[i] + 1$. If token is not idle, follow rule to release critical section.

Suzuki Kasami Algorithm

- To enter critical section:
 - Enter Critical Section if the token is received after sending REQUEST
- To release critical section:
 - set $LN[i] = RN_i[i]$
 - For every process j which is not in Q (in token), add process j to Q if $RN_i[j] = LN[j] + 1$
 - If Q is non-empty after the above, delete the first process from Q and send the token to that process

Suzuki Kasami Algorithm - Analysis

- **Message complexity:** zero if node holds the token already and token is idle, n otherwise
- **Synchronization delay:** zero (node has the token) or max. message delay (token is elsewhere)
- No starvation

